

# Agentic AI Architecture

A Deep Dive into Autonomous Systems, Design Patterns, and  
Implementation Frameworks

# Agenda

- \* Agentic AI Architecture
- \* Agentic Architecture Types
- \* Key Components of the Agentic AI Framework
  - \* Perception Module
  - \* Cognitive Module
  - \* Action Module
  - \* Learning Module
  - \* Collaboration Module
  - \* Security Module
- \* Agentic AI Design Patterns
  - \* Reflection Pattern
  - \* Tool Use Pattern
  - \* Planning Pattern
  - \* ReAct (Reasoning and Acting)
  - \* ReWOO (Reasoning WithOut Observation)
  - \* Multi Agent Pattern
  - \* Design Considerations

# Core Concepts

---

Understanding the paradigm shift from Automation to Autonomy

# The Evolution of AI Systems

## Traditional Automation

Systems follow rigid, pre-defined rules. If condition A is met, execute action B. These systems are deterministic, fragile to change, and require manual intervention for edge cases. They "do" exactly what they are told.

## Agentic AI

Systems possess agency. They are given a high-level goal and must autonomously determine the steps to achieve it. They perceive their environment, reason about actions, use tools, and adapt to feedback to solve complex problems.

# The Agentic Control Loop

## Perceive, Reason, Act

At the heart of every agent is a continuous control loop.

- The agent observes the state of the world (Perception),
- Processes this information against its goals and memory (Cognitive),
- Executes commands via tools (Action).
- This loop enables the system to interact dynamically with its environment, creating a feedback mechanism that allows for error correction and multi-step problem solving.

# Agentic Architecture Types



## Single Agent

A solitary autonomous entity with a unified memory and toolset. Best for linear, straightforward tasks where context switching is minimal and global context is required.



## Multi-Agent System (MAS)

A network of specialized agents working together. Each agent focuses on a distinct domain (e.g., Coder, Reviewer). Best for complex, multifaceted problems requiring diverse expertise.

# Key Components

---

Deconstructing the Agentic Framework

# 1. Perception Module

## Sensing the Environment

The perception module is the agent's interface with the world.

It converts external signals into a format the Cognitive Module can process.

- **Multimodal Inputs:** Text, Images, Audio, and Video.
- **Digital Signals:** API responses, database states, and system logs.
- **User Intent:** parsing complex natural language instructions.

## 2. Cognitive Module

### The Brain (LLM)

This is the core processing unit, typically a Large Language Model (LLM). It handles reasoning, planning, and decision-making.

It maintains **Memory** (Short-term context and Long-term vector stores) to retain conversation history and knowledge. It uses this context to decide which tool to use and *how* to use it.



# 3. Action & 4. Learning Modules

## Action Module

The "hands" of the agent. This module executes the decisions made by the brain. It includes:

- Tool Use: Calling external APIs (search, calculator)
- Effectors: Manipulating files, sending emails, or controlling robotic hardware.
- Output: Generating final responses

## Learning Module

How the agent improves over time. Unlike static scripts, agents can adapt.

- In-Context Learning: Adapting behavior based on prompt history.
- RAG (Retrieval): Fetching relevant knowledge to update its context window.
- Fine-tuning: Updating model weights for specific tasks.

# 5. Collaboration & 6. Security

## Collaboration Module

Essential for Multi-Agent Systems. It defines how agents talk to each other.

- **Orchestration:** A central manager delegates tasks.
- **Choreography:** Agents hand off tasks autonomously.
- **Shared Memory:** A blackboard for shared state.

## Security Module

The guardrails preventing rogue actions.

- **Input Validation:** Preventing prompt injection.
- **Action Permissibility:** Limiting API scope (e.g., read-only access).
- **Human-in-the-Loop:** Requiring approval for high-stakes actions.

# Design Patterns

---

Architectural blueprints for robust  
agents

# Pattern: Reflection

## Self-Correction Loop

Reflection allows an agent to critique its own output before finalizing it. This mimics human revision.

- ⟳ **Workflow:** Generate -> Critique -> Regenerate.
- ✓ **Benefit:** Significantly reduces hallucination and logical errors.
- </> **Use Case:** Code generation (writing code, running tests, fixing errors based on stderr).

# Patterns: Tool Use & Planning



## Tool Use

Extending the LLM's capability beyond text. The model outputs structured JSON to call functions (e.g., `get\_weather(location)`), bridging the gap between language and logic.



## Planning

Decomposing complex goals. Techniques like Chain of Thought (CoT) and Tree of Thoughts (ToT) break big problems into manageable, sequential steps.

# Reasoning Architectures

## ReAct

Reason + Act. The agent reasons about the state, takes an action, observes the output, and repeats.

Pros: Highly accurate, handles dynamic environments well.

Cons: High latency, sequential execution creates bottlenecks.

## ReWOO

Reason WithOut Observation. The agent generates a full plan *upfront* with placeholders for tool outputs, then executes tools in parallel.

Pros: Much faster, lower token cost.

Cons: Less adaptive if a tool fails unexpectedly.

# Multi-Agent Patterns

## From Solo to Squad

Complex problems often exceed the context window or reasoning capabilities of a single model. Multi-Agent Systems (MAS) solve this by distributing tasks.

- **Specialization:** Agents assume distinct personas (e.g., 'Researcher', 'Writer').
- **Orchestration:** A root agent manages the workflow and consolidates results.
- **Scalability:** Parallel execution reduces total processing time.

# Multi-Agent Collaboration Models

-  **Hierarchical (Boss/Worker):** A "Manager" agent breaks down the task and assigns sub-tasks to specialized "Worker" agents (e.g., Coder, Researcher). The Manager aggregates results.
-  **Sequential Handoffs:** A linear chain where the output of Agent A becomes the input for Agent B. Ideal for well-defined pipelines (e.g., Write -> Review -> Publish).
-  **Joint Collaboration:** Agents act as peers in a chat room, discussing the problem until a consensus is reached. Good for creative brainstorming.

# Design Considerations

---

Challenges in  
Production

# Implementation Challenges

## Latency

### Inference Time

Chained calls add up. Parallelization  
(ReWOO) is key.

## Cost

### Token Usage

Verbose reasoning loops can be  
expensive. Optimize prompts.

## Loops

### Error Recovery

Agents can get stuck in infinite retry  
loops. Set max-steps.